

What is a database?

→ Storing large amounts of data in a structured way

→ Allows for dynamic queries for data

→ It used to be about storing:

→ Corporate data

→ Payrolls, inventory, sales, customers, accounting, documents, ...

→ Banking Systems

→ Stock Exchanges

→ Airline Systems

→ Etc.



→ Today, databases are used in all fields:

→ Web backends:

→ Web search (google, bing, etc.)

→ Social Networks (Facebook, Instagram, etc.)

→ Blogs, Forums, etc.

→ Mobile applications

→ Data Warehouses

→ Big Data

→ AI

→ Etc.

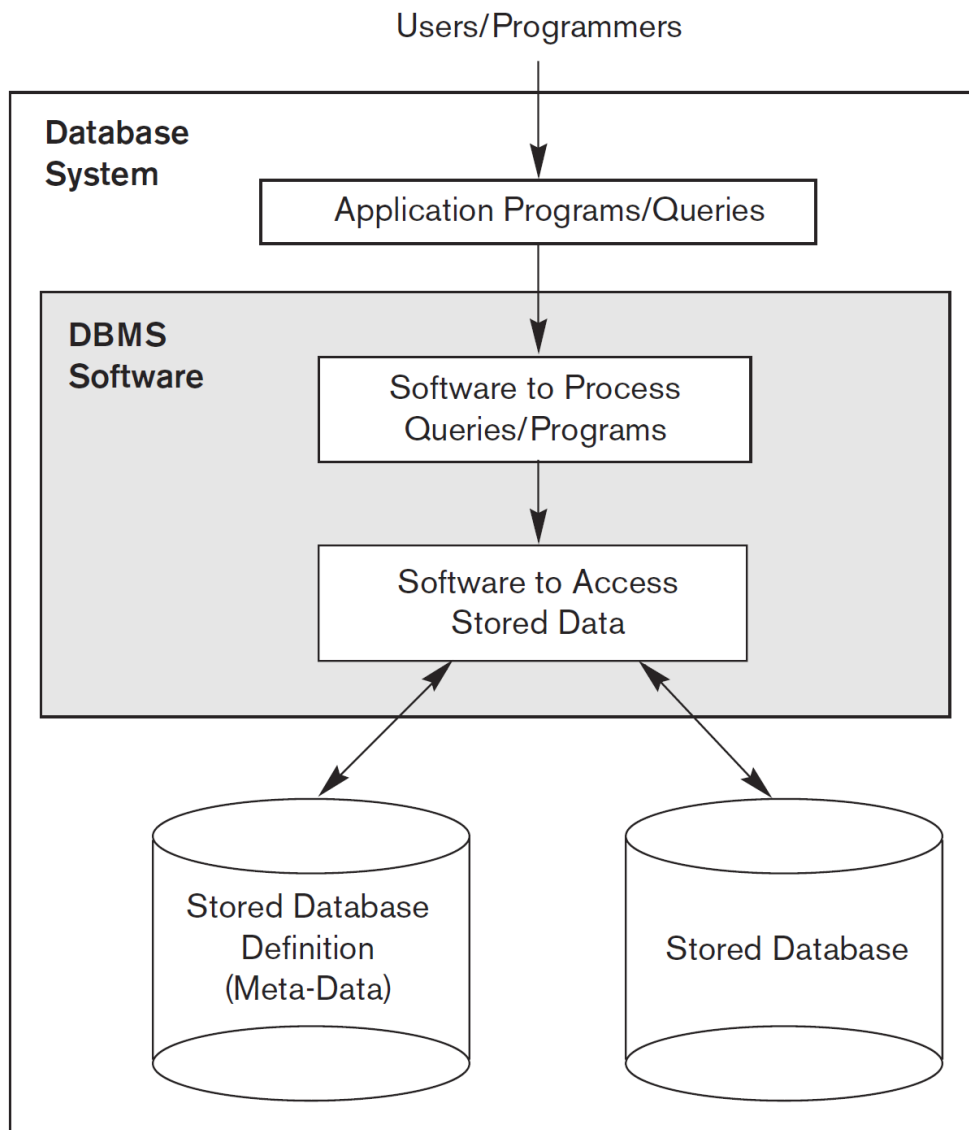


Why use a database?

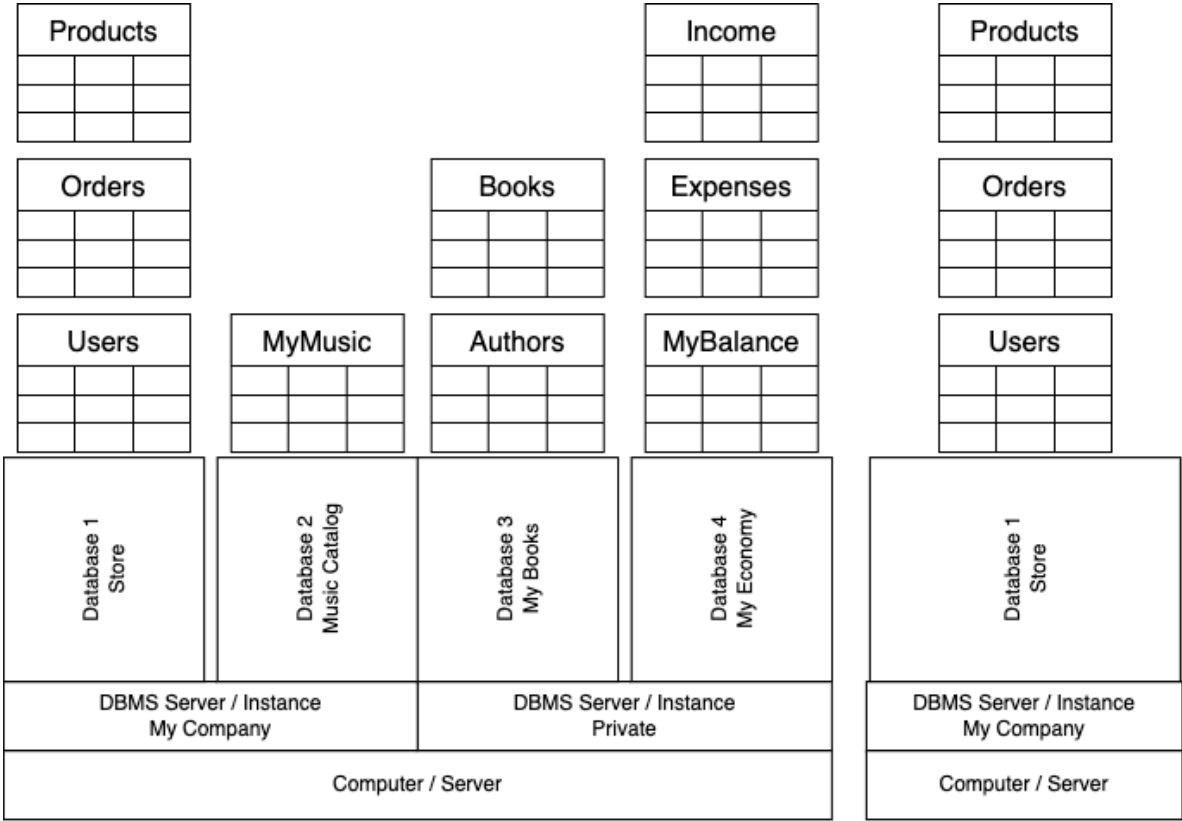
- Easy to use
- Flexible search
- Efficient
- Centralized storage
- Multi-user access
- Scalability
- Security and consistency

Definitions 1/4


- Database Management System (DBMS) – A system that enables users to create and maintain a database.
- Database – A collection of related data. A collection of random data is not a database. Often the word database also, incorrectly, refers to a DBMS.
- Data – Known facts that can be recorded that has implicit meaning.
- Mini-World / Universe of Discourse – a database that represents some aspect of the real world. (Mostly here as the book refers to it)
- Meta data – Data about data, or data that gives context to other data.
- Schema – A definition of table structures and their relationships.
- SQL – Structured Query Language – A language to extract data from a database.




Definitions 2/4



Definitions 3/4



Definitions 3/4



- Table – A collection of rows and columns that contains Cells.
- Row, Tuple, or Record – A collection of cells that together form a set of data that has a concrete Relation to each other. Shown as the horizontal line to the left.
- Column, Attribute, or Field – A collection of Cells that are all the same type. They are a part of multiple Rows.
- Cell – A specific Rows Column. Ergo the intersection.
- Relation – The relationship between data in a Row.
- Value – The information saved in the specific Cell.
- View / Result Set – The table that is created in memory and sent to the client as a result of a query.

Relational Databases

- Consists of multiple tables
- Tables relate to each other
- Uses primary and foreign keys to enforce the relationship constraints


Rank			DBMS	Database Model	Score		
Feb 2020	Jan 2020	Feb 2019			Feb 2020	Jan 2020	Feb 2019
1.	1.	1.	Oracle +	Relational, Multi-model i	1344.75	-1.93	+80.73
2.	2.	2.	MySQL +	Relational, Multi-model i	1267.65	-7.00	+100.36
3.	3.	3.	Microsoft SQL Server +	Relational, Multi-model i	1093.75	-4.80	+53.69
4.	4.	4.	PostgreSQL +	Relational, Multi-model i	506.94	-0.25	+33.38
5.	5.	5.	IBM Db2 +	Relational, Multi-model i	165.55	-3.15	-13.87
6.	6.	6.	Microsoft Access	Relational	128.06	-0.52	-15.96
7.	7.	7.	SQLite +	Relational	123.36	+1.22	-2.81
8.	8.	8.	MariaDB +	Relational, Multi-model i	87.34	-0.11	+3.91
9.	9.	↑ 10.	Hive +	Relational	83.53	-0.71	+11.25
10.	10.	↓ 9.	Teradata +	Relational, Multi-model i	76.81	-1.48	+0.84
11.	↑ 12.	↑ 12.	SAP HANA +	Relational, Multi-model i	54.97	+0.28	-1.58
12.	↓ 11.	↓ 11.	FileMaker	Relational	54.88	-0.23	-2.91
13.	13.	13.	SAP Adaptive Server	Relational	52.73	-1.86	-3.02
14.	14.	14.	Microsoft Azure SQL Database	Relational, Multi-model i	31.41	+3.20	+4.28
15.	15.	↑ 20.	Google BigQuery +	Relational	27.56	+0.81	+8.81

NoSQL Databases

→ NoSQL means **Not only SQL**


→ NoSQL covers multiple types of databases

Data Model	Example Databases
Key-Value (“Key-Value Databases,” p. 81)	BerkeleyDB
	LevelDB
	Memcached
	Project Voldemort
	Redis
	<i>Riak</i>
Document (“Document Databases,” p. 89)	CouchDB
	<i>MongoDB</i>
	OrientDB
	RavenDB
	Terrastore
Column-Family (“Column-Family Stores,” p. 99)	Amazon SimpleDB
	<i>Cassandra</i>
	HBase
	Hypertable
Graph (“Graph Databases,” p. 111)	FlockDB
	HyperGraphDB
	Infinite Graph
	<i>Neo4J</i>
	OrientDB



What is a relational database?

- Has a Schema that defines the Tables
- Uses multiple tables to store data
- Uses the notion of Primary Keys and Foreign keys to relate table content to each other.
- Uses SQL which makes CRUD (Create, Read, Update, Delete) operations on Schemas, Tables, and Rows.



PostgreSQL

PostgreSQL

- This course will use PostgreSQL (AKA Postgres)
- However, this is **NOT** a PostgreSQL course!

- PostgreSQL is:
 - A Relational Database
 - Cross-Platform (Windows, Mac OS, Linux, BSD, Solaris)
 - Licensed under the PostgreSQL License, a liberal Open Source license, similar to the BSD or MIT licenses.
 - Open Source (<https://github.com/postgres/postgres>)
- Documentation and usage sites:
 - <https://www.postgresql.org/docs/12/index.html>
 - <https://www.postgresqltutorial.com/>

```
-- creating the initial accounts table
CREATE TABLE account(
  user_id serial PRIMARY KEY,
  username VARCHAR (50) UNIQUE NOT NULL,
  password VARCHAR (50) NOT NULL,
  email VARCHAR (355) UNIQUE NOT NULL,
  created_on TIMESTAMP NOT NULL,
  last_login TIMESTAMP
);

-- inserting two users
INSERT INTO account (username, password, email, created_on)
VALUES ('John', 'myPassW0rd', 'john@acme.com', NOW());

INSERT INTO account (username, password, email, created_on)
VALUES ('Anne', 'myPassW0rd', 'anne@acme.com', NOW());

-- querying for all rows in the account table
SELECT * FROM account;

-- updating the password of the user Anne
UPDATE account SET password = 'newPassW0rd' WHERE username = 'Anne';
```

What is SQL?

- Acronym for: **Structured Query Language**.
- Pronounced either Sequel, or SQL.
- Allows for retrieval of data from a database.
- A query results in a result set, which is structured as a table with rows and columns.

Tables and relationships

STUDENT

Name	Student_number	Class	Major
Smith	17	1	CS
Brown	8	2	CS

COURSE

Course_name	Course_number	Credit_hours	Department
Intro to Computer Science	CS1310	4	CS
Data Structures	CS3320	4	CS
Discrete Mathematics	MATH2410	3	MATH
Database	CS3380	3	CS

PREREQUISITE

Course_number	Prerequisite_number
CS3380	CS3320
CS3380	MATH2410
CS3320	CS1310

SECTION

Section_identifier	Course_number	Semester	Year	Instructor
85	MATH2410	Fall	07	King
92	CS1310	Fall	07	Anderson
102	CS3320	Spring	08	Knuth
112	MATH2410	Fall	08	Chang
119	CS1310	Fall	08	Anderson
135	CS3380	Fall	08	Stone

GRADE_REPORT

Student_number	Section_identifier	Grade
17	112	B
17	119	C
8	85	A
8	92	A
8	102	B
8	135	A

Creating a Table

```
CREATE TABLE account(  
  user_id serial PRIMARY KEY,  
  username VARCHAR (50) UNIQUE NOT NULL,  
  password VARCHAR (50) NOT NULL,  
  email VARCHAR (355) UNIQUE NOT NULL,  
  created_on TIMESTAMP NOT NULL,  
  last_login TIMESTAMP  
);
```

	<div>user_id</div> <div>[PK] integer</div>	<div>username</div> <div>character varying (50)</div>	<div>password</div> <div>character varying (50)</div>	<div>email</div> <div>character varying (355)</div>	<div>created_on</div> <div>timestamp without time zone</div>	<div>last_login</div> <div>timestamp without time zone</div>

Inserts (CRUD)

```
INSERT INTO account (username, password, email, created_on)
VALUES ('John', 'myPassW0rd', 'john@acme.com', NOW());
```

```
INSERT INTO account (username, password, email, created_on)
VALUES ('Anne', 'myPassW0rd', 'anne@acme.com', NOW());
```

	user_id [PK] integer	username character varying (50)	password character varying (50)	email character varying (355)	created_on timestamp without time zone	last_login timestamp without time zone
1	4	John	myPassW0rd	john@acme.com	2020-02-06 11:43:44.158522	[null]
2	5	Anne	myPassW0rd	anne@acme.com	2020-02-06 11:43:44.158522	[null]

Selects (CRUD)

```
SELECT * FROM account;

SELECT username, created_on FROM account WHERE email = 'john@acme.com';

SELECT username, created_on FROM account WHERE email LIKE '%anne%';
```

	<div>user_id</div> <div>[PK] integer</div>	<div>username</div> <div>character varying (50)</div>	<div>password</div> <div>character varying (50)</div>	<div>email</div> <div>character varying (355)</div>	<div>created_on</div> <div>timestamp without time zone</div>	<div>last_login</div> <div>timestamp without time zone</div>
1	1	John	myPassW0rd	john@acme.com	2020-02-06 11:41:11.729203	[null]

Updates (CRUD)

```
UPDATE account SET password = 'newPassW0rd' WHERE username = 'Anne';
```

	<div>user_id</div> <div>[PK] integer</div>	<div>username</div> <div>character varying (50)</div>	<div>password</div> <div>character varying (50)</div>	<div>email</div> <div>character varying (355)</div>	<div>created_on</div> <div>timestamp without time zone</div>	<div>last_login</div> <div>timestamp without time zone</div>
1	1	John	myPassW0rd	john@acme.com	2020-02-06 11:41:11.729203	[null]
2	2	Anne	newPassW0rd	anne@acme.com	2020-02-06 11:42:13.27506	[null]

Delete (CRUD)

```
DELETE FROM account WHERE email = 'john@acme.com';
```

	<div>user_id</div> <div>[PK] integer</div>	<div>username</div> <div>character varying (50)</div>	<div>password</div> <div>character varying (50)</div>	<div>email</div> <div>character varying (355)</div>	<div>created_on</div> <div>timestamp without time zone</div>	<div>last_login</div> <div>timestamp without time zone</div>
1	5	Anne	myPassW0rd	anne@acme.com	2020-02-06 11:43:44.158522	[null]

BEWARE of doing this: `DELETE FROM account;`

Constraints and Data types in SQL

```
CREATE TABLE account(  
    user_id serial PRIMARY KEY,  
    username VARCHAR (50) UNIQUE NOT NULL,  
    password VARCHAR (50) NOT NULL,  
    email VARCHAR (355) UNIQUE NOT NULL,  
    created_on TIMESTAMP NOT NULL,  
    last_login TIMESTAMP  
);
```

Constraints – Abbreviated, more later.

- PRIMARY KEY – The unique identifier for the current row, which is always NOT NULL
- FOREIGN KEY – A key that refers to a primary key in a different table, signifying their relationship to each other
- UNIQUE – Two cells in this Column cannot be the same
- NOT NULL – This attribute HAS to be specified

PostgreSQL Data Types 1/3

→ Null – Value Missing

→ Boolean – 1 bit

→ Converts Boolean values e.g., 1, yes, y, t, true are converted to true, and 0, no, n false, f are converted to false.

→ Character types

→ CHAR(n) – Fixed length text. Unused space is padded with space characters.

→ VARCHAR(n) – Variable length string, as in “store up to”.

→ TEXT – Large size text lengths such as book texts.

PostgreSQL Data Types 2/3

→ Numeric types

→ Integer

- SMALLINT – 2 byte, ranges from -32.768 to 32.767 (2 byte = 2 x 8 bit = 256 x 256 = 65.536)
- INT – 4 byte integer, range from -2,147,483,648 to 2,147,483,647
- SERIAL – Same as INT, but used for auto incrementing.

→ Floating-point

- FLOAT(n) – floating point number with at the precision of n, and a maximum of 8 bytes. (n as in numbers after the “,”)
- REAL – A floating point number. 0.001, 0.00000001, etc.

→ Temporal types

- DATE – dates only
- TIME – time of day values
- TIMESTAMP – both date and time values
- INTERVAL – periods of time

PostgreSQL Data Types 3/3

- UUID for storing Universally Unique Identifiers
 - Array for storing array strings, numbers, etc.
 - JSON stores JSON data
 - hstore stores key-value pair
 - Special types
 - box, line, point, lseg, polygon, inet, macaddr.
- See more here: <https://www.postgresql.org/docs/12/datatype.html>