


NoSQL Databases

→ NoSQL means **Not Only SQL**

→ NoSQL covers multiple types of databases

Data Model	Example Databases
Key-Value (“Key-Value Databases,” p. 81)	BerkeleyDB LevelDB Memcached Project Voldemort Redis <i>Riak</i>
Document (“Document Databases,” p. 89)	CouchDB <i>MongoDB</i> OrientDB RavenDB Terrastore
Column-Family (“Column-Family Stores,” p. 99)	Amazon SimpleDB <i>Cassandra</i> HBase Hypertable
Graph (“Graph Databases,” p. 111)	FlockDB HyperGraphDB Infinite Graph <i>Neo4J</i> OrientDB

Source: NoSQL Distilled, by P. Sadalage and M. Fowler



Why NoSQL?

- Fits well to many data types and application areas
- Dynamic Schema! (Bad and Good)
- Easy to scale
 - Data Size increases (Big Data)
- Often easy to replicate
- High Performance
- Less powerful query languages
- Versioning

A comment on database sizes (Relational DBs)

					id	Symbol	1	Price	Timestamp
<input type="checkbox"/>		Ret			16777216	BTCUSDT		23494.4	1677554256
<input type="checkbox"/>		Ret			33554432	BTCUSDT		27888	1679411028
<input type="checkbox"/>		Ret			50331648	BTCUSDT		29282.3	1682851218
<input type="checkbox"/>		Ret			65536	BTCUSDT		20976.3	0
<input type="checkbox"/>		Ret			16842752	BTCUSDT		23407.1	1677560723
<input type="checkbox"/>		Ret			33619968	BTCUSDT		28347.6	1679418043
<input type="checkbox"/>		Ret			50397184	BTCUSDT		29203.9	1682858445
<input type="checkbox"/>		Ret			131072	BTCUSDT		21120.8	0
<input type="checkbox"/>		Ret			16908288	BTCUSDT		23393.2	1677567301
<input type="checkbox"/>		Ret			33685504	BTCUSDT		28070	1679424984
<input type="checkbox"/>		Ret			50462720	BTCUSDT		29255	1682865889

✓ Viser rækkerne 0 - 24 (5977768 i alt, Forespørgsel tog 24.4428 sekunder.)

About 6.000.000 rows

A comment on database sizes (Relational DBs)

✓ MySQL returnerede ingen data (fx ingen rækker). (Forespørgsel tog 18.7534 sekunder.)

```
SELECT * from price_logs WHERE Symbol LIKE "BTCUSD*" LIMIT 20;
```

Takes 18,7 seconds to run

✓ MySQL returnerede ingen data (fx ingen rækker). (Forespørgsel tog 87.5783 sekunder.)

```
CREATE INDEX symbol_index on price_logs(Symbol);
```

Create an index to optimize the 'like' clause, which takes 87,5783 seconds to create in memory.

✓ MySQL returnerede ingen data (fx ingen rækker). (Forespørgsel tog 0.0002 sekunder.)

```
SELECT * from price_logs WHERE Symbol LIKE "BTCUSD*" LIMIT 20;
```

After optimization, the query takes 0,0002 seconds to run (or 0.2 milliseconds)

Key-Value




Graph DB




Four NOSQL Categories

BigTable



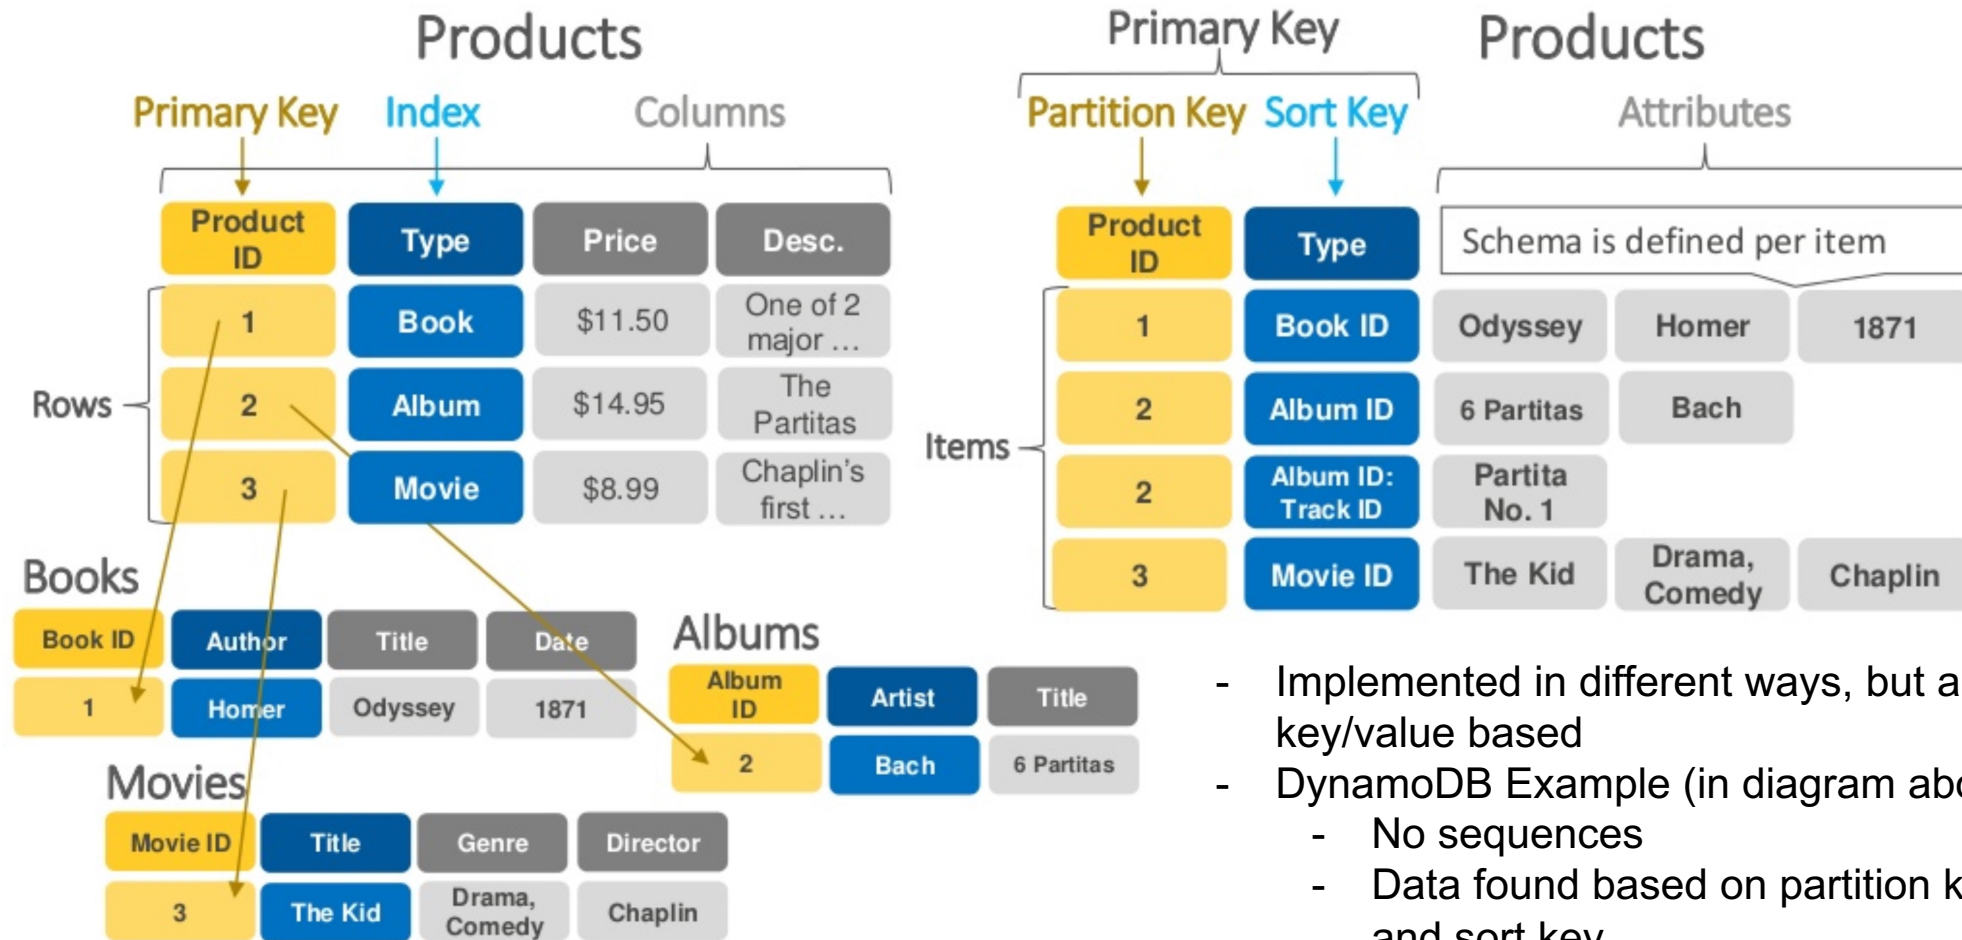
Document



What types exist?


- Document-based NoSQL systems
- NoSQL key-value stores
- Column based or wide column-based NoSQL systems
- Graph based NoSQL systems
- Hybrid NoSQL systems
 - Object databases
 - XML databases

Relational vs. Key Value

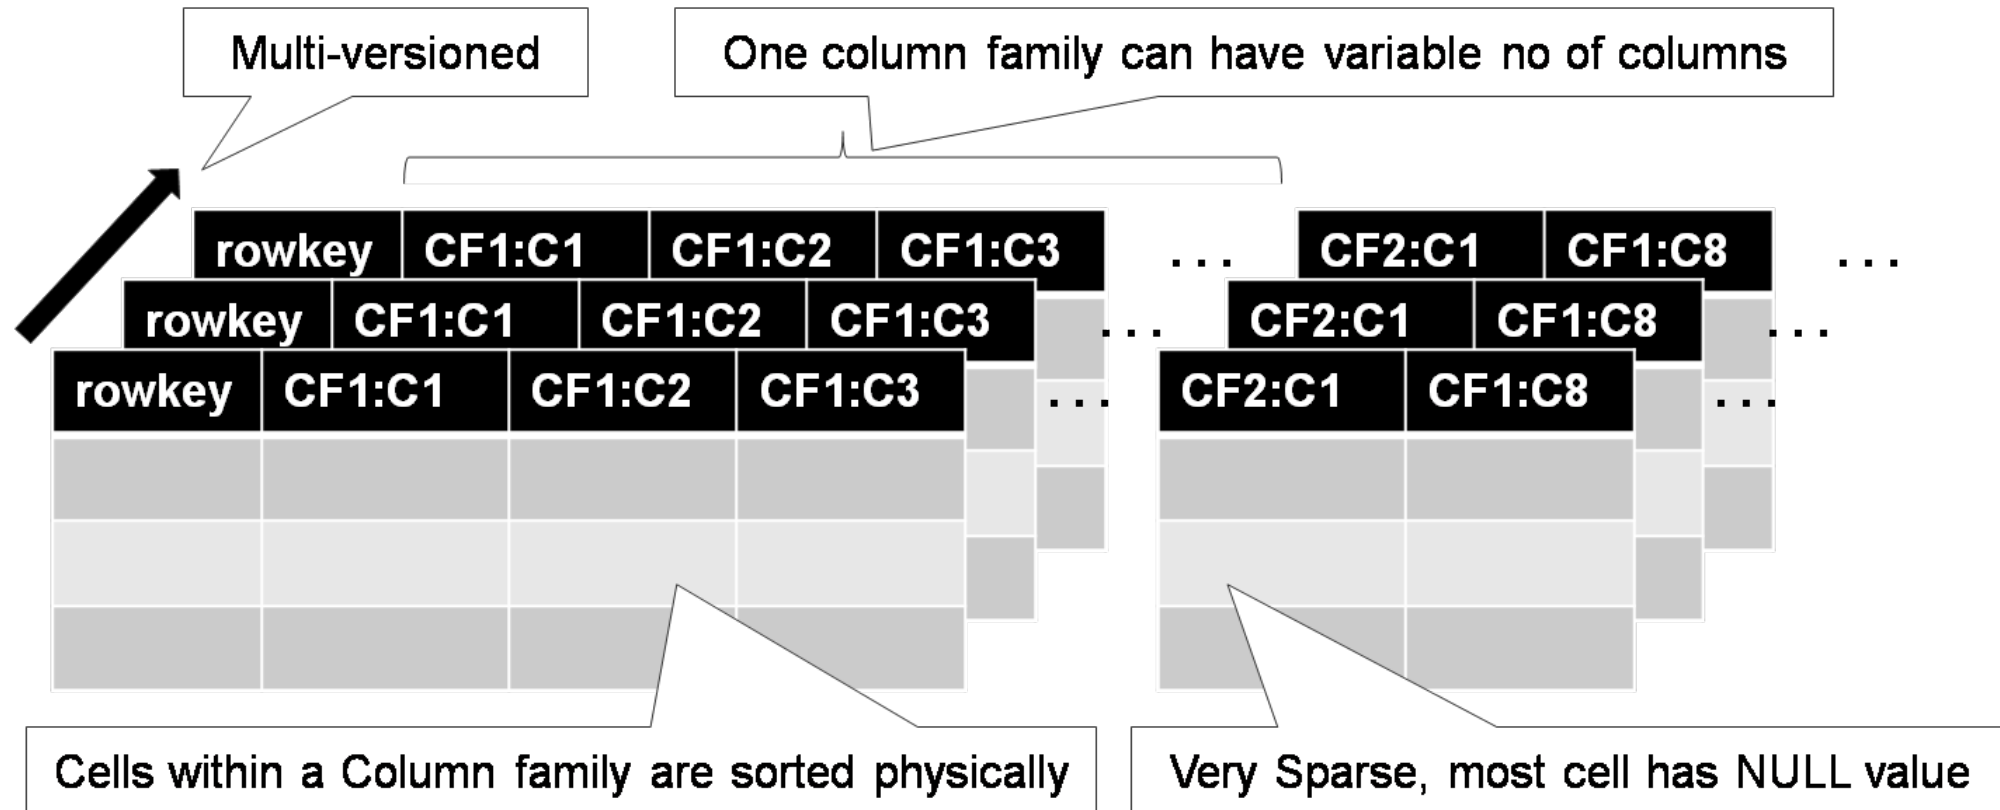


- Implemented in different ways, but all key/value based
- DynamoDB Example (in diagram above)
 - No sequences
 - Data found based on partition key and sort key

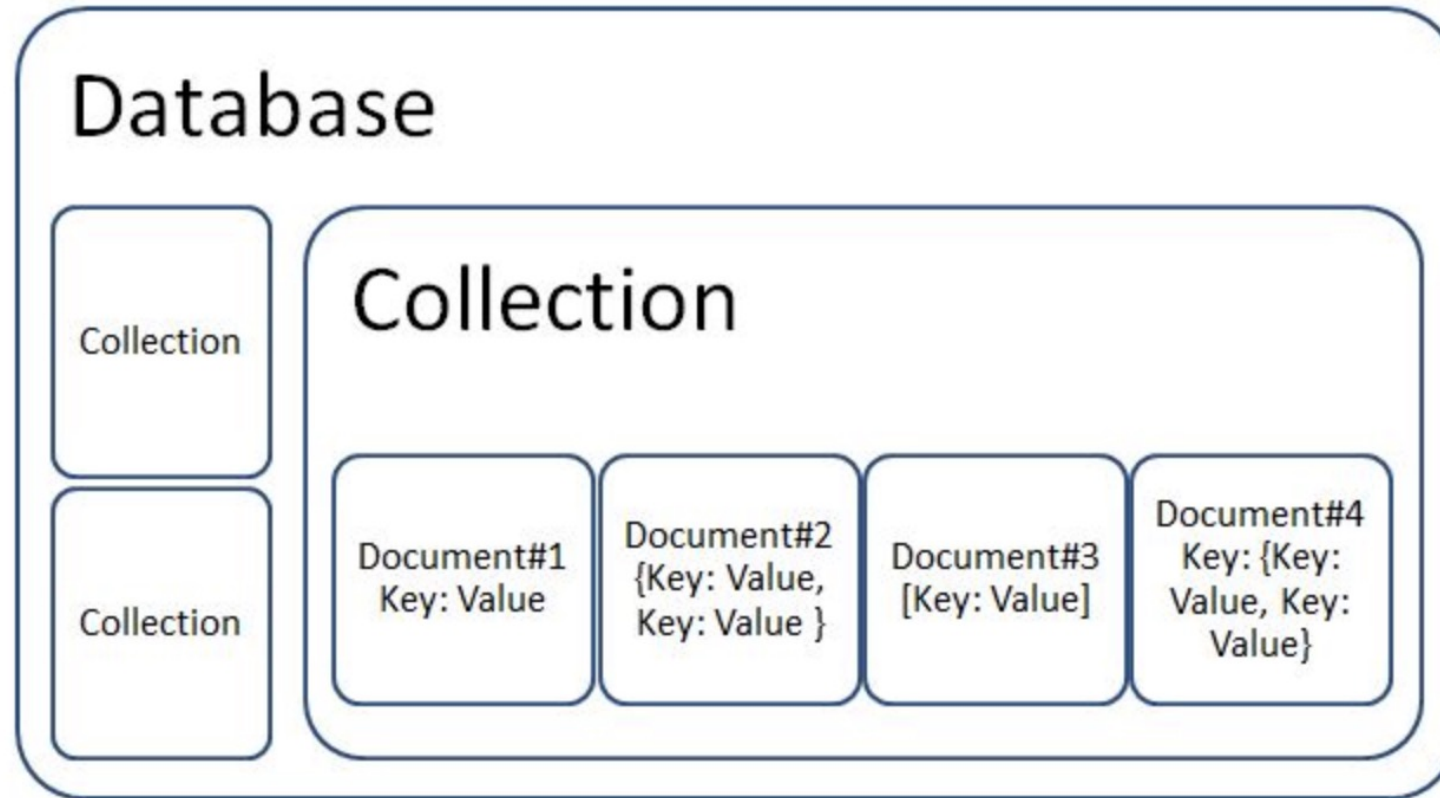
Graph Based



Big Table




Document based



Horizontal vs Vertical Scaling


Vertical Scaling

(Increase size of instance (RAM , CPU etc.))



Horizontal Scaling

(Add more instances)



→ Vertical

→ Increasingly Expensive to get larger hardware

→ Requires no code changes

→ Horizontal

→ Less expensive hardware

→ Allows for large-scale systems

→ Need to change software for a distributed architecture


→ Leads to more complex code

→ Need for Load Balancing


→ Do you know examples of these?

High Availability 1/3

Active/Passive Cluster





Active/Passive Cluster - Failover




High Availability 2/3



Active/Active Cluster




Active/Active Cluster - Failure




Single Point of Failure



No Single Point of Failure








Partitioning (Sharding)

- Splitting up the data between multiple servers
- Together they represent the entire collection
- Allows for better throughput as calculations and search are shared
- Makes sense for Big Data Applications

Partitioning (Sharding)

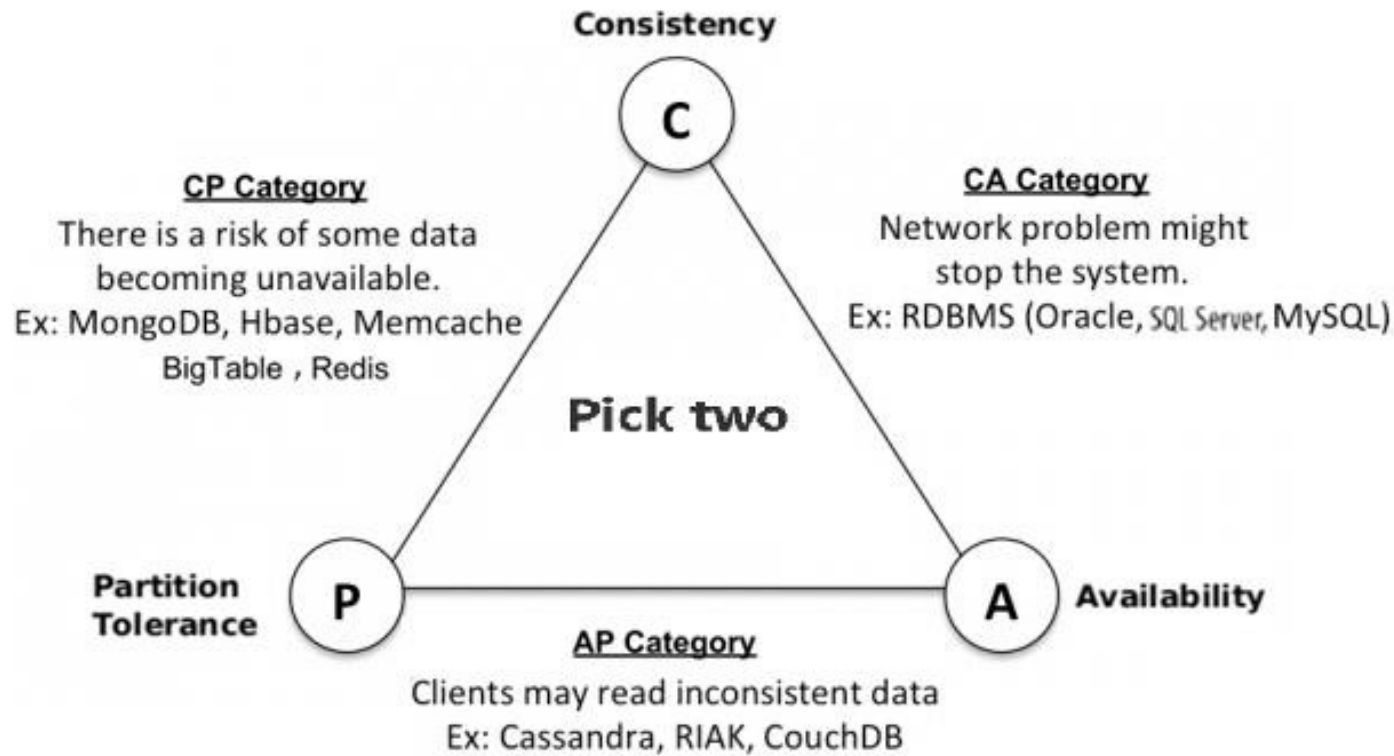


Vertical



Horizontal

The CAP Theorem – Distributed Databases




Consistency: Every read receives the most recent write or an error. All nodes see the same data at all times.

Availability: System is always operational despite individual server status. Every Request receives a (non-error) response, without the guarantee that it contains the most recent write

Partition Tolerance: The system continues to operate despite of partition failure (a node dies with part of the data)


Consistency



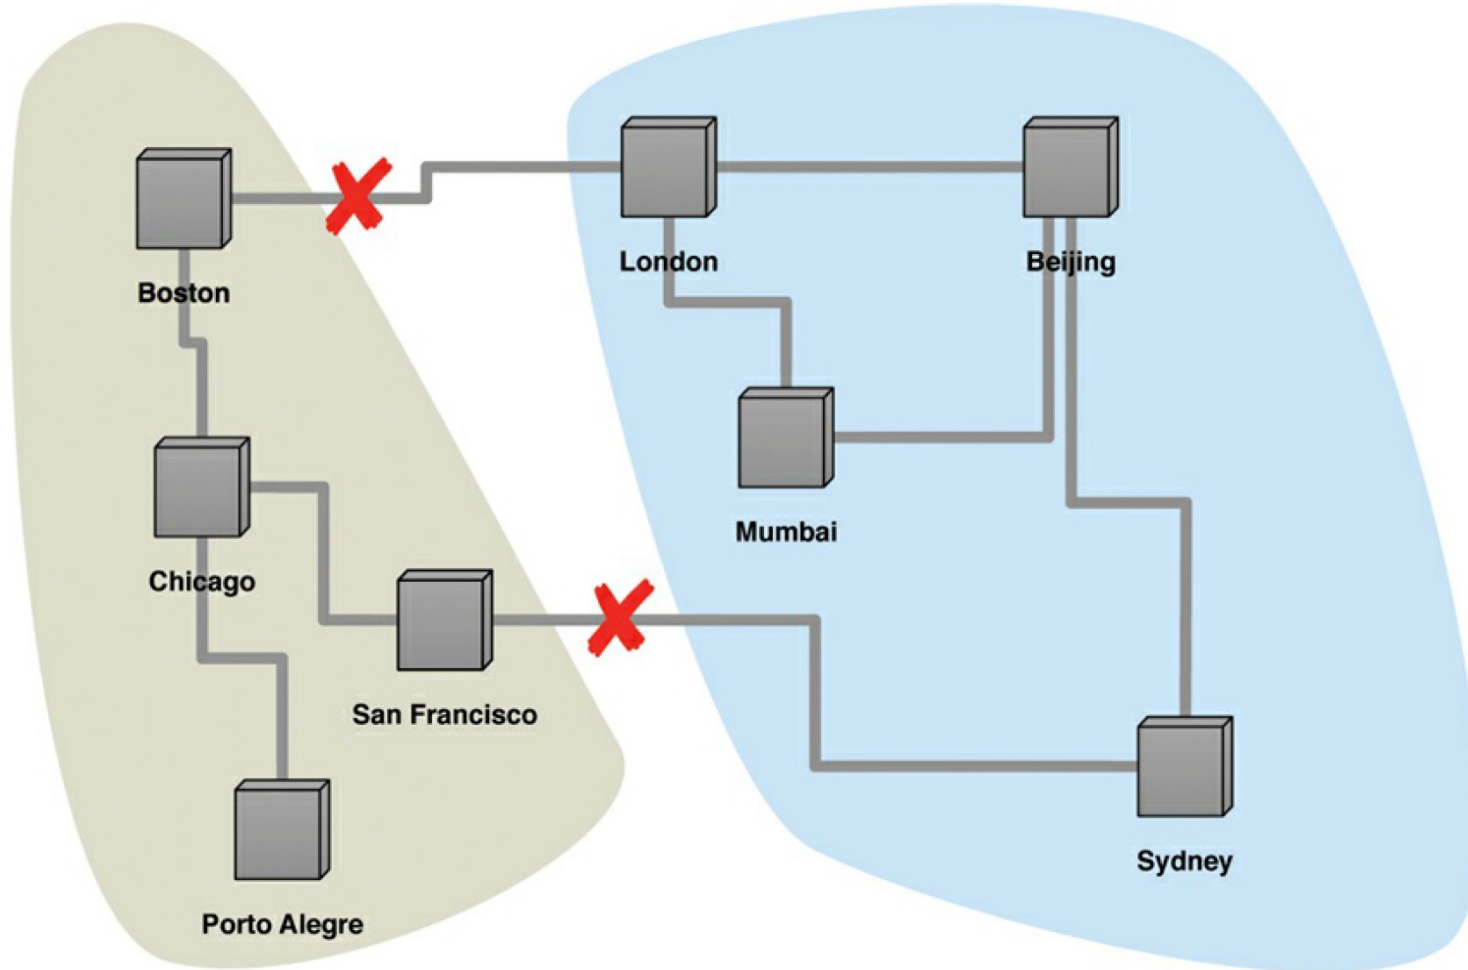
Availability



Partition tolerance



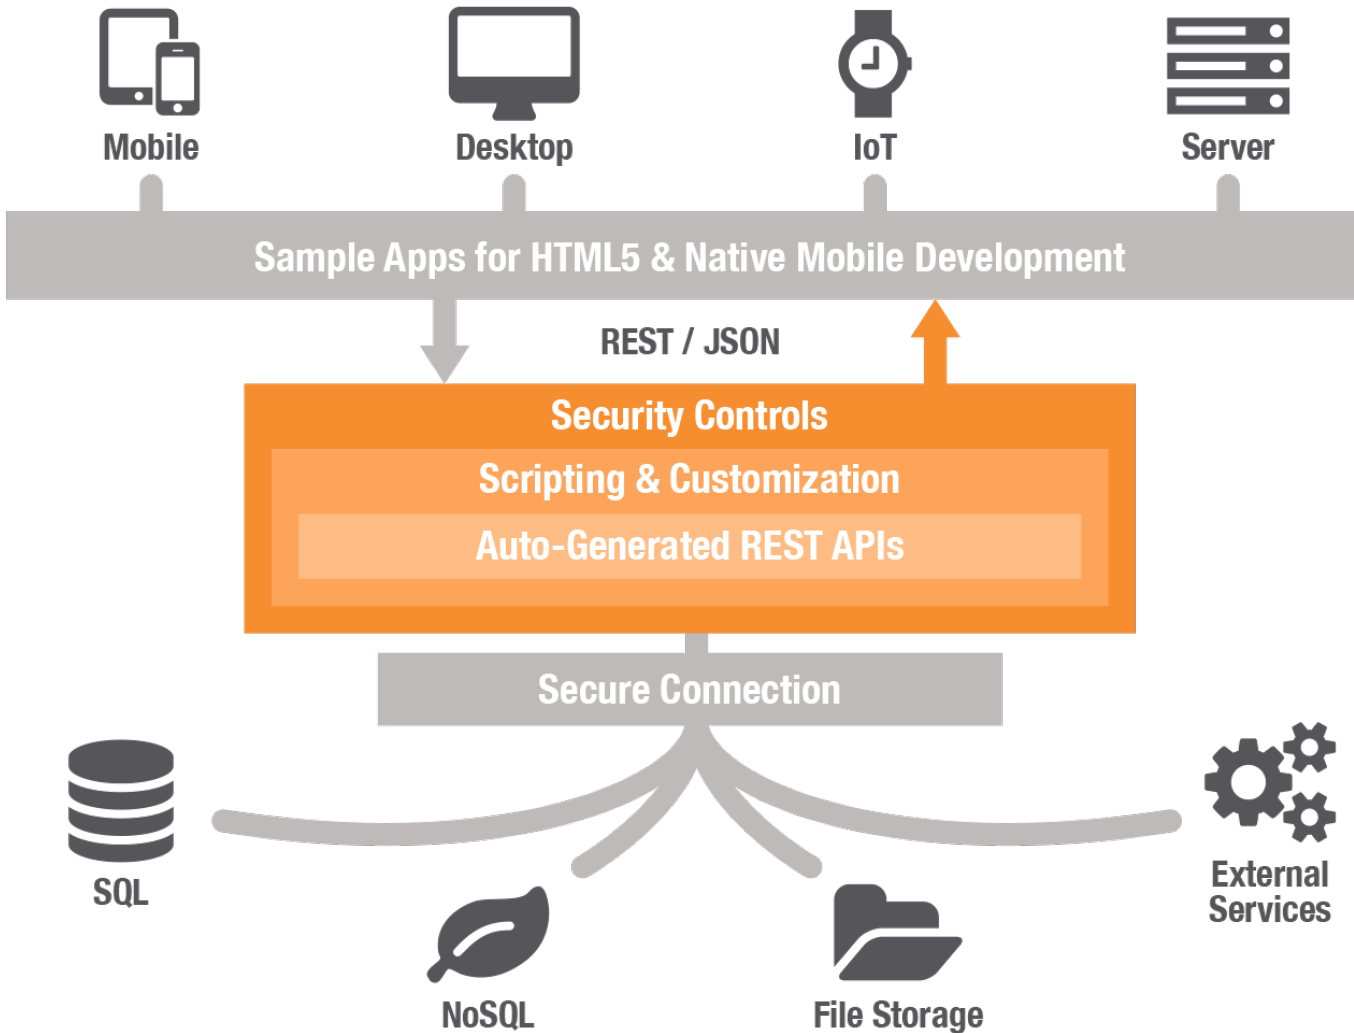
A CAP Theorem example





What is JSON?

- JSON stands for “JavaScript Object Notation”,
- It’s used to store and exchange data as an alternative solution for XML.
- Easy to parse, and easy to read and write for a human.
- Based on JavaScript Object Literals, but it’s a text format.
- JSON is language independent; means you can use parse and generate JSON data in other programming languages.



How is JSON used?

- Websites communicating with the backend
- Mobile applications communicating with the backend
- Temporary storage format of objects
- ...

Object Declaration

```
{  
  "id": 2456,  
  "name": "John Doe",  
  "cpr": "111111-1111",  
  "married": true  
}
```


- An object starts with { and ends with }.
- Inside is a collection of Key/Values separated by comma
 - Be aware it is a syntax error to have a comma after the last key/value
- Each key is a string, so it uses the double quotes - "key"
- A : mark separates a key and a value
- Each value can be of several types. In the example we have number, string, and boolean.

Array Declaration

```
[  
    "Odense",  
    "Aarhus",  
    "San Francisco"]
```

- Arrays are ordered lists, and starts from 0
- Initiated with [and ends with]
- Values are separated by comma
- An array can also hold objects

JSON Values



- string; "string"
- number; 1
- boolean; true or false
- null; null/left out
- an object; {"key": "value"}
- an array; [1, "2", null, true]

Putting it together

```
{  
  "id": 2456,  
  "name": "John Doe",  
  "cpr": "111111-1111",  
  "married": true,  
  "relationships": [  
    {  
      "name": "Jane Doe",  
      "relationshipType": "Wife"  
    },  
    {  
      "name": "Danny Doe",  
      "relationshipType": "Son"  
    }  
  ],  
  "livedIn": ["Odense", "Aarhus"]  
}
```